

HIGH PERFORMANCE COMPUTING REPORT ON
“PARALLEL IMPLEMENTATION AND EVALUATION OF
QUICK SORT USING OPENMPI “

SUBMITTED TO THE SAVTRIBAI PHULE PUNE UNIVERSITY, PUNE
IN THE PARTIAL FULFILLMENT OF THE REQUIREMENTS
FOR THE AWARD OF THE DEGREE

OF

FINAL YEAR (COMPUTER ENGINEERING)

BY

Sumit Kumar	Roll no: C41152
Ashwin Bettawar	Roll no: C41105
Nandini Agrawal	Roll no: C41101

Under the guidance of
Prof. A. A. Bhise



Sinhgad Institutes

DEPARTMENT OF COMPUTER ENGINEERING

STES'S SMT. KASHIBAI NAVALE COLLEGE OF ENGINEERING
VADGAON BK, OFF SINHGAD ROAD, PUNE 411041
SAVTRIBAI PHULE PUNE UNIVERSITY

2022 – 23



Sinhgad Institutes

DEPARTMENT OF COMPUTER ENGINEERING

STES'S SMT. KASHIBAI NAVALE COLLEGE OF ENGINEERING

VADGAON BK, OFF SINHGAD ROAD, PUNE 411041

CERTIFICATE

This is to certify that the project report entitles

**“PARALLEL IMPLEMENTATION AND EVALUATION OF QUICK SORT
USING OPENMPI”**

Submitted by

Sumit Kumar	Roll no: C41152
Ashwin Bettawar	Roll no: C41105
Nandini Agrawal	Roll no: C41101

is a bonafide work carried out by them under the supervision of **Prof. A. A. Bhise** and it is submitted towards the partial fulfillment of the requirement of SAVTRIBAI PHULE PUNE UNIVERSITY, Pune for the award of the degree (Computer Engineering)

(Prof. A. A. Bhise)

Guide

Department of Computer Engineering

(Prof. R. H. Borhade)

Head,

Department of Computer Engineering

(Dr. A. V. Deshpande)

Principal,

Smt. Kashibai Navale College of Engineering Pune – 41

Place: Pune

Date: / /

ACKNOWLEDGEMENT

We would like to express our very great appreciation to **Prof. A. A. Bhise** for his valuable and constructive suggestions during the planning and development of this research work. His willingness to give his time so generously has been very much appreciated. His patient guidance, enthusiastic encouragement, useful critiques and also his assistance in keeping our progress on schedule throughout our semester will be acknowledged.

Sumit Kumar	Roll no: C41152
Ashwin Bettawar	Roll no: C41105
Nandini Agrawal	Roll no: C41101

ABSTRACT

Sorting is used in human activities and devices like personal computers, smart phones, and it continues to play a crucial role in the development of recent technologies. The Quicksort algorithm has been known as one of the fastest and most efficient sorting algorithms. It was invented by C.A.R Hoare in 1961 and is using the divide-and-conquer strategy for solving problems [3]. Its partitioning aspects make Quicksort amenable to parallelization using task parallelism. MPI is a message passing interface library allowing parallel computing by sending codes to multiple processors, and can therefore be easily used on most multi-core computers available today. The main aim of this study is to implement the Quicksort algorithm using the Open MPI library and therefore compare the sequential with the parallel implementation. The entire study will be divided in many sections: related works, theory of experiment, algorithm and implementation of Quicksort using open MPI, the experimental setup and results, problems followed by the conclusion and future works.

TABLE OF CONTENTS

1. INTRODUCTION	7
2. THEORY	8
3. PROPOSED ALGORITHM	10
4. IMPLEMENTATION	11
5. SPECIFICATIONS	15
5.1 HARDWARE SPECIFICATIONS	15
5.2 SOFTWARE SPECIFICATIONS	15
6. RESULTS	16
7. VISUALISATION	18
8. ANALYSIS	20
9. CONCLUSION	21
10. REFERENCES	22

1. INTRODUCTION

Sorting is used in human activities and devices like personal computers, smart phones, and it continues to play a crucial role in the development of recent technologies. The Quick Sort algorithm has been known as one of the fastest and most efficient sorting algorithms. It was invented by C.A.R Hoare in 1961 and is using the divide-and-conquer strategy for solving problems [3]. Its partitioning aspects make Quick Sort amenable to parallelization using task parallelism. MPI is a message passing interface library allowing parallel computing by sending codes to multiple processors, and can therefore be easily used on most multi-core computers available today. The main aim of this study is to implement the Quick Sort algorithm using the Open MPI library and therefore compare the sequential with the parallel implementation. The entire study will be divided in many sections: related works, theory of experiment, algorithm and implementation of Quicksort using open MPI, the experimental setup and results, problems followed by the conclusion and future works.

2. THEORY

Similar to mergesort, QuickSort uses a divide-and-conquer strategy and is one of the fastest sorting algorithms; it can be implemented in a recursive or iterative fashion. The divide and conquer is a general algorithm design paradigm and key steps of this strategy can be summarized as follows:

- **Divide:** divide the input data set S into disjoint subsets $S_1, S_2, S_3 \dots S_k$.
- **Recursion:** solve the sub-problems associated with $S_1, S_2, S_3 \dots S_k$
- **Conquer:** combine the solutions for $S_1, S_2, S_3 \dots S_k$ into a solution for S .
- **Base case:** the base case for the recursion is generally sub-problems of size 0 or 1.

Many studies [2] have revealed that in order to sort N items; it will take the QuickSort an average running time of $O(N \log N)$. The worst-case running time for QuickSort will occur when the pivot is a unique minimum or maximum element, and as stated in [2] the worst-case running time for QuickSort on N items is $O(N^2)$. These different running times can be influenced by the inputs distribution (uniform, sorted or semi-sorted, unsorted, duplicates) and the choice of the pivot element. Here is a simple Pseudocode of the QuickSort algorithm adapted from Wikipedia

```
function quicksort(array)
    less, equal, greater := three empty arrays
    if length(array) > 1
        pivot := select any element of array
        for each x in array
            if x < pivot then add x to less
            if x = pivot then add x to equal
            if x > pivot then add x to greater
        quicksort(less)
        quicksort(greater)
    array := concatenate(less, equal, greater)
```

As shown in the above Pseudocode, the following is the basic working mechanism of Quicksort:

- Choose any element of the array to be the pivot
- Divide all other elements (except the pivot) into two partitions
 - All elements less than the pivot must be in the first partition
 - All elements greater than the pivot must be in the second partition
 - Use recursion to sort both partitions
 - Join the first sorted partition, the pivot and the second partition
 - The output is a sorted array.

We have made use of Open MPI as the backbone library for parallelizing the QuickSort algorithm. In fact learning message passing interface (MPI) allows us to strengthen our fundamentals knowledge on parallel programming, given that MPI is lower level than equivalent libraries (OpenMP). As simple as its name means, the basic idea behind MPI is that messages can be passed or exchanged among different processes in order to perform a given task. An illustration can be a communication and coordination by a master process which split a huge task into chunks and share them to its slave processes. Open MPI is developed and maintained by a consortium of academic, research and industry partners; it combines the expertise, technologies and resources all across the high performance computing community [11]. As elaborated in [4], MPI has two types of communication routines: point-to-point communication routines and collective communication routines. Collective routines as explained in the implementation section have been used in this study.

3. PROPOSED ALGORITHM

In general, the overall algorithm used here to perform QuickSort with MPI works as followed:

- i. Start and initialize MPI.
- ii. Under the root process MASTER, get inputs:
 - a. Read the list of numbers L from an input file.
 - b. Initialize the main array globaldata with L.
 - c. Start the timer.
- iii. Divide the input size SIZE by the number of participating processes npes to get each chunk size localsize.
- iv. Distribute globaldata proportionally to all processes:
 - a. From MASTER scatter globaldata to all processes.
 - b. Each process receives in a sub data localdata.
- v. Each process locally sorts its localdata of size localsize.
- vi. Master gathers all sorted localdata by other processes in globaldata.
 - a. Gather each sorted localdata.
 - b. Free localdata.
- vii. Under MASTER perform a final sort of globaldata.
 - a. Final sort of globaldata.
 - b. Stop the timer.
 - c. Write the output to file.
 - d. Sequentially check that globaldata is properly and correctly sorted.
 - e. Free globaldata.
- viii. Finalize MPI.

4. IMPLEMENTATION

In order to implement the above algorithm using C programming, we have made use of a few MPI collective routine operations. Therefore after initializing MPI with **MPI_Init**, the size and rank are obtained respectively using **MPI_Comm_size** and **MPI_Comm_rank**. The beginning wall time is received from **MPI_Wtime** and the array containing inputs is distributed proportionally to the size to all participating processes with **MPI_Scatter** by the root process which collect them again after they are sorted using **MPI_Gather**. Finally the ending wall time is retrieved again from **MPI_Wtime** and the MPI terminate calling **MPI_Finalize**. In the following part, each section of our proposed algorithm is illustrated with a code snippet taken from the implementation source code.

- i. Start and initialize MPI.

```
MPI_Init(&argc, &argv);
MPI_Comm_size(MPI_COMM_WORLD, &npes);
MPI_Comm_rank(MPI_COMM_WORLD, &rank);
```

- ii. Under the root process *MASTER*, get inputs:

```
if (rank == MASTER) {
    printf("SIZE IS %lld", SIZE);
    globaldata = malloc (SIZE * sizeof(long long) );
    if (globaldata == NULL) {
        printf ("\n\n globaldata Memory Allocation Failed !  \n\n ");
        exit(EXIT_FAILURE);
    }
    inp = fopen (argv[1], "r");          /* Open file for reading */
    if (inp == NULL) {
        printf ("\n\n inp Memory Allocation Failed !  \n\n ");
        exit(EXIT_FAILURE);
    }
    printf ("\n\nInput Data \n\n ");
    for (i=0; i<SIZE; i++) {
        retscan = fscanf (inp, "%lld \t", &tmp);
        globaldata[i] = tmp;
    }
}
```

Starting the timer.

```
if (rank == MASTER)
{
    t_start = MPI_Wtime ();
}
```

- i. Divide the input size **SIZE** by the number of participating processes **npes** to get each chunk size **localsize**.

```
/*Getting the size to be used by each process */
if (SIZE < npes) {
    printf ("\n\n SIZE is less than the number of process! \n\n ");
    exit (EXIT_FAILURE);
}
localsize = SIZE/npes;
```

- ii. Distribute **globaldata** proportionally to all processes:

```
/*Scatter the integers to each number of processes (npes) */
MPI_Scatter (globaldata, localsize, MPI_LONG_LONG, localdata,
            localsize, MPI_LONG_LONG, MASTER, MPI_COMM_WORLD);
```

- i. Each process locally sorts its **localdata** of size **localsize**.

```
/* Perform local sort on each sub data by each process */
quickSortRecursive (localdata,0, localsize-1);
```

- ii. **Master** gathers all sorted **localdata** by other processes in **globaldata**.

```
/* Merge locally sorted data of each process by MASTER to globaldata */
MPI_Gather (localdata, localsize, MPI_LONG_LONG , globaldata,
            localsize, MPI_LONG_LONG, MASTER, MPI_COMM_WORLD);
free (localdata);
```

- iii. Under **MASTER** perform a final sort of *globaldata*.

```
if (rank == MASTER) {  
    /* Final sorting */  
    quickSortRecursive (globaldata, 0, SIZE-1);  
}
```

Stop the timer

```
/* End wall time */  
t_end = MPI_Wtime ();
```

Write information to in the output file

```
/* Opening output file to write sorted data */  
out = fopen (argv[2], "w");  
if (out == NULL) {  
    printf ("\n\n out Memory Allocation Failed ! \n\n ");  
    exit (EXIT_FAILURE);  
}  
/* Write information to output file */  
fprintf (out, "Recursively Sorted Data : ");  
fprintf (out, "\n\nInput size : %lld\t", SIZE);  
fprintf (out, "\n\nNber processes : %d\t", npes);  
fprintf (out, "\n\nWall time : %7.4f\t", t_end - t_start);  
printf ("\n\nWall time : %7.4f\t", t_end - t_start);  
fprintf (out, "\n\n");  
for (i = 0; i<SIZE; i++) {  
    fprintf (out, " %lld \t", globaldata[i]);  
}  
fclose (out); /* closing the file */
```

Checking that the final *globaldata* array is properly sorted.

```
/* checking if the final globaldata content is properly sorted */  
sortCheckers ( SIZE, globaldata );
```

Free the allocated memory

```

if (rank == MASTER) {
    free (globaldata);
}
/* free the allocated memory */

```

ix. Finalize MPI.

```

/* MPI_Finalize Terminates MPI execution environment */
MPI_Finalize ();

```

The two versions of QuickSort algorithm have been implemented; however even though they have almost the same implementation using similar functions, the recursion in the recursive part has been replaced by a stack in order to make it iterative. Their function signatures are presented in the following part:

- Recursive QuickSort

```

void quickSortRecursive (long long [], long long, long long);
long long partition (long long [], long long, long long);
void swap (long long [], long long, long long);
void sortCheckers (long long, long long []);
long long getSize (char str[]);

```

- Iterative QuickSort

```

void quickSortIterative (long long [], long long, long long);
long long partition (long long [], long long, long long);
void swap (long long [], long long, long long);
void sortCheckers (long long, long long []);
long long getSize (char str[]);

```

The input file necessary to run this program can be generated using an input generator where we specify the size and the filename (*input_generator.c*), then compile and run with the following instructions:

5. SPECIFICATIONS

5.1 HARDWARE SPECIFICATIONS

- Computer : ProBook-4440s
- Memory : 3,8 GiB
- Processor : Intel® Core™ i5-3210M CPU @ 2.50GHz × 4
- Graphics : Intel® Ivybridge Mobile x86/MMX/SSE2
- OS type : 32-bit
- Disk : 40,1 GB

5.2 SOFTWARE SPECIFICATIONS

- Operating system : Ubuntu 14.04 LTS
- Open MPI : 1.10.1 released
- Compilers : gcc version 4.8.4
- Microsoft Excel 2007 for plotting

6. RESULTS




The following table presents the different recorded data. In the first column we have the experiment number (No.); the second column is the number of participating processes (# process), the third column is the input data size applied to QuickSort. Finally the last two columns represent respective the execution wall time of the iterative and recursive version of parallel QuickSort.



No.	# process	Input size	Iterative wall time	Recursive wall time
1	1	20	0.0000	0.0000
	2		0.0001	0.0000
	5		0.0002	0.0004
	10		0.0004	0.0005
	20		0.0013	0.0032
2	1	100	0.0000	0.0000
	2		0.0001	0.0001
	5		0.0002	0.0004
	10		0.0003	0.0005
	20		0.0016	0.0020
3	1	1000	0.0011	0.0012
	2		0.0003	0.0003
	5		0.0004	0.0004
	10		0.0007	0.0008
	20		0.0014	0.0016
4	1	10000	0.0849	0.0860
	2		0.0030	0.0030
	5		0.0031	0.0030
	10		0.0038	0.0035
	20		0.0035	0.0043
5	1	100000	8.2165	8.5484
	2		0.0393	0.0383
	5		0.0333	0.0325
	10		0.0418	0.0488
	20		0.0446	0.0475
6	1	1000000	835.8316	2098.7

	2		0.4786	0.4471
	5		0.3718	0.3590
	10		0.3646	0.3445
	20		0.4104	0.3751

7. VISUALISATION

Different charts comparing the iterative QuickSort and the recursive QuickSort for different number of processes and various input sizes are presented in this section. The X-axis represents the number of processes and the Y-axis represents the execution wall time in seconds.





Here is a sample output file format we have used along this experiment.

```
Iteratively Sorted Data :
Input size : 100
Nber processes : 10
Wall time      : 0.0003
2  3  5  5  8  11  11  12  13  13
```

8. ANALYSIS

In Figure 01 and Figure 02, the sorting is faster with a single process and keeps slowing as long as the number of processes increases. However in Figure 03 the execution time drastically dropped from a single process to its fastest execution time where the two processes are running in parallel, then it starts slowing again when we increases the number of processes. Finally in Figure 04 and Figure 05 the iterative and recursive implementations have almost the same execution time, the sorting becomes faster with more processes, with less variations. Figure 06 having one million numbers as input size. Here on single process, the sorting is the slowest of this experiment and even stop on recursion. The sorting becomes again faster with the number of process increases. On these different charts we can clearly observe that in general the iterative QuickSort is faster than the recursion version. On sequential execution with a single process, the sorting is faster with small input and slows down as long as the input size goes up. However on parallel execution with more than one process, the execution time decreases when the input size and the number of process increases. We noticed sometimes an unusual behavior of MPI execution time that keeps changing after each execution. We have taken only the first execution time to minimize this variation and obtain a more consistent execution time.

9. CONCLUSION

To conclude this project, we have successfully implemented the sequential QuickSort algorithm both the recursive and iterative version using the C programming language. Then we have done a parallel implementation with the help of Open MPI library. Through this project we have explored different aspects of MPI and the QuickSort algorithm as well as their respective limitations. This study has revealed that in general the sequential QuickSort is faster on small inputs while the parallel QuickSort excels on large inputs. The study also shows that in general the iterative version of QuickSort perform slightly better than the recursive version. The study reveals some unusual behavior of the execution time that unexpectedly varies. However due to the limited time we could not overcome all the difficulties and limitations, the next section opens an eventual scope for future improvements.

10. REFERENCES

1. Rosetta code, available at http://rosettacode.org/wiki/Sorting_algorithms/Quicksort , December 2015.
2. P. Tsigas and Y. Zhang. A Simple, Fast Parallel Implementation of Quicksort and its Performance Evaluation on SUN Enterprise 10000. In Proceedings of the Eleventh Euromicro Conference on Parallel, Distributed and Network-Based Processing (Euro-PDP'03), 2003, IEEE.
3. HOARE, C. A. R. 1961. Algorithm 64: Quicksort. Commun. ACM 4, 7, 321.
4. Blaise Barney, Lawrence Livermore National Laboratory, available at <https://computing.llnl.gov/tutorials/mpi/>, December 2015
5. Rer. nat. Siegm. Groß, available at http://www2.hsfulda.de/~gross/cygwin/cygwin_2.0/inst_conf_cygwin_2.0.x_x11r7.x.pdf , December 2015
6. D. G. Martinez, S. R. Lumley, Installation of Open MPI parallel and distributed programming.
7. P. Kataria, Parallel quicksort implementation using MPI and Pthreads, December 2008
8. P. Perera, Parallel Quicksort using MPI and Performance Analysis
9. V. Prifti, R. Bala, I. Tafa, D. Saatciu and J. Fejzaj. The Time Profit Obtained by Parallelization of Quicksort Algorithm Used for Numerical Sorting. In Science and Information Conference 2015, June 28-30, 2015, IEEE.
10. V. Gite, UNIX Command Line Tools For MS-Windows XP / Vista / 7 Operating Systems, available at <http://www.cyberciti.biz/faq/unix-command-line-utilities-for-windows/> , December 2015
11. Open MPI: Open Source High Performance Computing, available at <https://www.open-mpi.org/>, December 20